

## 1. Data Collection

## 2. Data Exploration

- Correlation
- Distribution
- Feature engineering

## 3. Data Processing

- Numeric value (missing value)
  - Interpolation
  - Scaling problem
- Non-Numeric value (Categorical variable)
  - Drop
  - Label Encoding
  - One-Hot Encoding

## 4. Train and Evaluating

- Supervised Learning
- Unsupervised Learning
- Semi-supervised Learning

## 5. Fine-Tune

- Optimal hyperparameter

### Skill:

- Sklearn, numpy, pandas
- python

### ML:

- Algorithm
- Cost function (utility function)
- Evaluation

## Linear Regression – linear datasets

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

$$y = h_{\theta}(\mathbf{x}) = \theta \cdot \mathbf{x}$$

**Essentiality** : weight \* features+ bias

**Performance measure:**

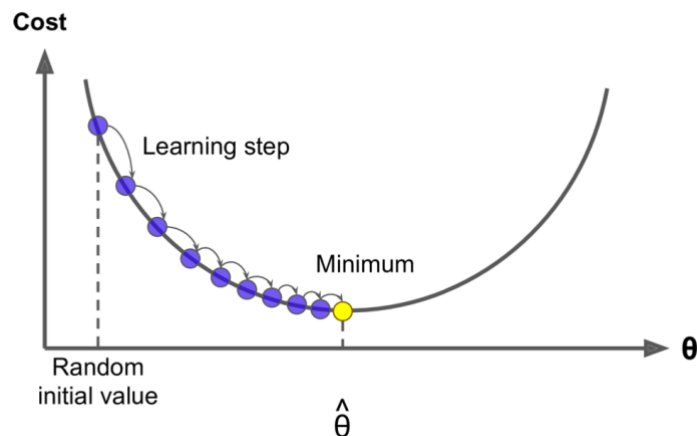
- MSE

$$MSE(X, h_{\theta}) = \frac{1}{m} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)})^2$$

- RMSE(Root Mean Square Error)
- R2

### Two way to train:

- a direct “closed-form” equation
  - the model parameters that minimize **the cost function** over the training set
  - Normal Equation
    - $\hat{\theta} = (X^T X^{-1}) X^T y$
    - SVD (Singular Value Decomposition)
- an iterative optimization approach (Gradient Descent)
  - tweak parameters iteratively in order to minimize a cost function
    - measures the local gradient of the error function with regards to the parameter vector  $\theta$ ,
    - goes in the direction of descending gradient
    - the gradient is zero = reach a minimum (global or local)

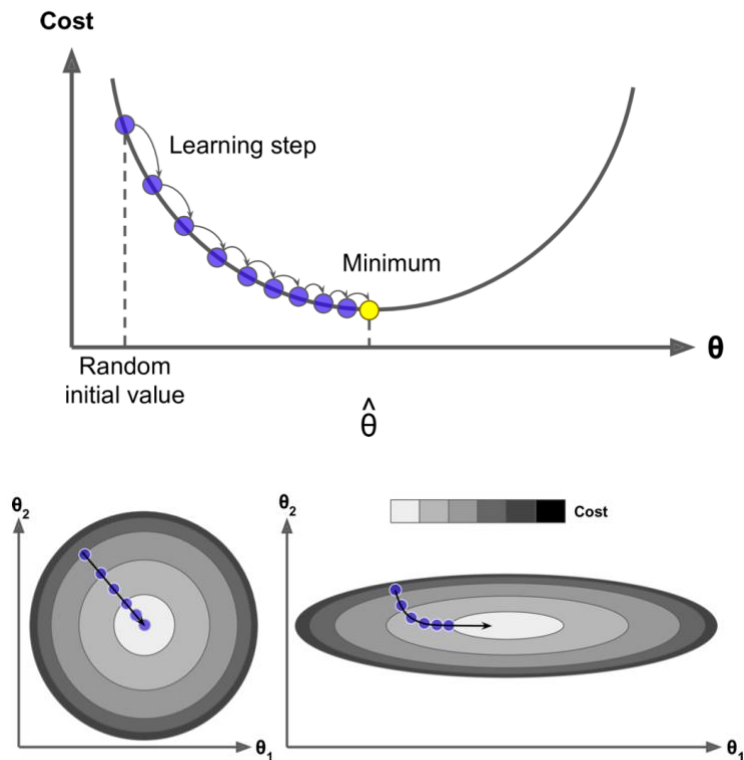


$$\nabla_{\theta} \text{MSE}(\theta) = \begin{pmatrix} \frac{\partial}{\partial \theta_0} \text{MSE}(\theta) \\ \frac{\partial}{\partial \theta_1} \text{MSE}(\theta) \\ \vdots \\ \frac{\partial}{\partial \theta_n} \text{MSE}(\theta) \end{pmatrix} = \frac{2}{m} \mathbf{X}^T (\mathbf{X}\theta - \mathbf{y})$$

$$\theta^{(\text{next step})} = \theta - \eta \nabla_{\theta} \text{MSE}(\theta)$$

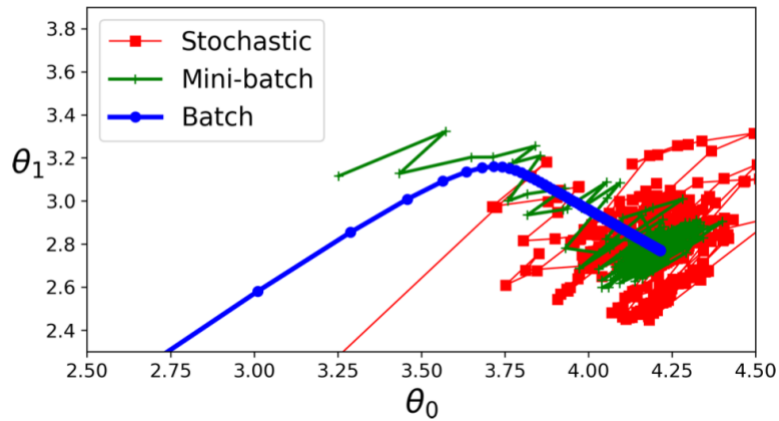
- **notice**
  - learning rate – how far a step will take
    - gradient – the direction at which the variable varies fastest
    - use grid search to find a good learning rate
      - tolerance – almost converge (minimum)
  - features scale
    - normalization

- standardization
- zero-centered

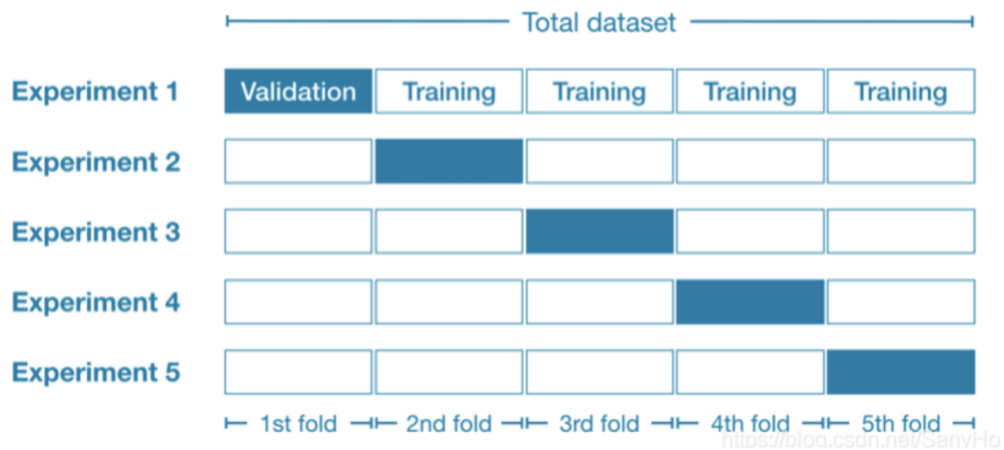


- Batch GD, Mini-batch GD, Stochastic GD
  - Batch GD -- the whole training set to compute the gradients at every step
  - Stochastic GD -- a random instance in the training set at every step
  - Mini-batch GD -- small random sets of instances

Algorithm	Large $m$	Out-of-core support	Large $n$	Hyperparams	Scaling required	Scikit-Learn
Normal Equation	Fast	No	Slow	0	No	n/a
SVD	Fast	No	Slow	0	No	LinearRegression
Batch GD	Slow	No	Fast	2	Yes	SGDRegressor
Stochastic GD	Fast	Yes	Fast	$\geq 2$	Yes	SGDRegressor
Mini-batch GD	Fast	Yes	Fast	$\geq 2$	Yes	SGDRegressor



## Cross-validation

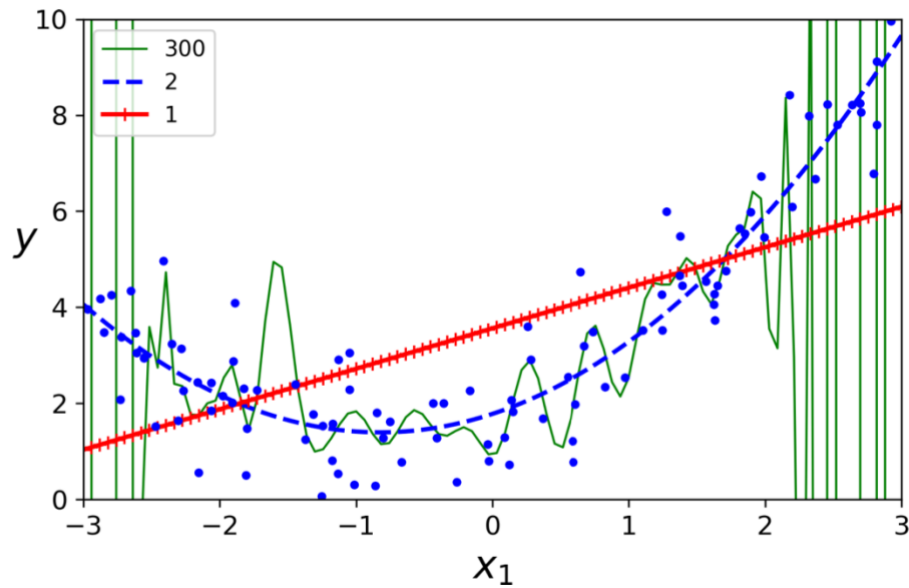


## Polynomial Regression – non-linear datasets

**Essence:** add power of each feature as new features, then train a linear model on the augmented data set.

### Drawback:

- Prone to be overfitting the training data
  - Learning curves
  - Regularization techniques



**Tradeoff** : bias, variance and irreducible error(noisiness)

**Fine-Tune** : search the optimal hyperparameter of the models

- **Grid Search**: evaluate **all the listed combinations** of hyperparameter values, using cross-validation
- **Randomized Search**: evaluate a given number of **random combinations** by selecting a **random value** for each hyperparameter at every iteration

**Learning Curves**: evaluate the model's generalization ability

- Underfitting : both curves reach a plateau, close and fairly high
- Overfitting : gap between the curves, error lower

**Regularized Linear Models**:

- the fewer degrees of freedom it has, the harder it will be for it to overfit the data
- Ridge Regression, Lasso Regression and Elastic Net
  - Ridge Regression
    - **cost function** :  $J(\theta) = MSE(\theta) + \alpha \frac{1}{2} \sum_{i=1}^n \theta_i^2$
    - **regularized term**( $L_2$ ):  $\alpha \frac{1}{2} \sum_{i=1}^n \theta_i^2 = \frac{1}{2} (||w||_2)^2$
    - bias, variance and irreducible error(noisiness)
      - $\alpha$  – reduced variance, increase bias
    - Sensitive to the scale of the input features (StandardScaler)
  - Lasso Regression
    - **cost function** :  $J(\theta) = MSE(\theta) + \alpha \sum_{i=1}^n |\theta_i|$
    - regularized term ( $L_1$ ) :  $\alpha \sum_{i=1}^n |\theta_i|$

- tend to completely eliminate the weights of the least important features
- performs feature selection and outputs a *sparse model*
- Elastic Net
  - **cost function:**  $J(\theta) = MSE(\theta) + r\alpha \sum_{i=1}^n |\theta_i| + \frac{1-r}{2} \alpha \sum_{i=1}^n \theta_i^2$
  - A middle ground between Ridge Regression and Lasso
- Plain Linear Regression < Ridge(regularized) < Lasso(feature selection, but erratically when number of features greater or some features strongly correlated) < Elastic Net
- 

**Early Stopping :** prevent the model from overfitting, get the best model

